

---

# DMA

— Direct Memory Access —

---

# Types of I/O Devices

## 1. Character devices

- E.g., keyboard, mouse
- Tend to have lower data rates
- Values (e.g., bytes) are transferred individually
- Per-value interrupts are appropriate


# Types of I/O Devices

## 1. Block devices

- E.g., Drive (mass storage device), video device, PCM sound device
- Tend to have higher data rates
- Entire blocks of values are transferred as a unit
- **Note:** Interrupts are not going away (when will the IRQ happen)?
- Why are per-value interrupts a problem?


# Approaches to Data Transfer

- Data is coming/going between the I/O interface and device
- As this happens, data also needs to be transferred from/to the I/O interface on the computer side (our focus in these slides)




# Approaches to Data Transfer

- E.g., once data has arrived from an input device ...
  - How is it obtained by the CPU?



# Approaches to Data Transfer

- E.g., once data has arrived from an input device ...
- How is it obtained by the CPU?
- How is it transferred to main memory for later use?



# Approaches to I/O

- There are three main approaches to I/O
  1. Programmed I/O
  2. Interrupt Driven I/O
  3. DMA (Direct Memory Access) I/O
- **Note:** Approaches to I/O are not mutually exclusive
  - More than one can be used in a single system

# Approaches to I/O - Programmed I/O

- Main (non-ISR) algorithm manages I/O directly
- Continually polls interface (no IRQs)
  - Transfers data when interface ready
- Requires least-sophisticated hardware
- Has highest software overhead
  - Example: Busy wait
  - Risk of over/under-run



# Approaches to I/O - Interrupts

- Install ISR that will be automatically CPU-invoked upon IRQ (if not masked)
  - Freeing up the logic of the main algorithm while increasing responsiveness
- ISR transfers data between interface and main memory (e.g., to/from a buffer)
  - The main algorithm only interacts with data in main memory
- Requires more sophisticated hardware
- Has much lower software overhead (even with ISR-vectoring overhead)

# Approaches to I/O - Interrupts

- Are there any drawbacks to interrupt-driven I/O?
  - It's better than polling, but the CPU is still involved in every I/O read/write operation. An ISR must be invoked for each one, which can be expensive – even if the CPU is just copying data from an I/O register to memory or vice versa!
  - This is too slow for high-speed devices which transfer large blocks of data at a time,
    - Example: hard drives
  - DMA (direct memory access) solves this problem by allowing direct I/O interface to memory transfers, and vice versa

# Approaches to I/O - DMA

- DMA (direct memory access) solves this problem by allowing direct I/O interface to memory transfers, and vice versa
  - Can go from device to memory **or** memory to device
  - Requires more sophisticated hardware
- A DMA controller manages the transfer, and becomes bus master as necessary to carry out the transfer
- Note that the CPU is still involved, in that it must initiate the DMA transfer. Additionally, it is typically notified of transfer completion via an interrupt

# DMA - Main Issues

- I/O controller must be provided with:
  - The operation to perform
  - A pointer to the main memory buffer
  - The size of the block
- I/O controller must maintain additional registers:
  - Index register (pointer into main memory)
  - Count register (number of bytes transferred to/from main memory)
- I/O controller must be able to drive the bus!
  - CPU is continually using the bus for transfers to/from main memory!
  - Simpler I/O interfaces are connected to the system bus via a DMA controller

# DMA Controller Registers

- A typical DMA controller has, minimally:
  - A control register (e.g. can control whether operation is read or write)
  - A status register
  - An address register
  - A data count register
  - 1-more data registers for buffering the data to be transferred

# Steps to Initiate DMA transfer

- The Software:
  1. Software prepares a main memory buffer
  2. Software requests an I/O operation
  3. Software then leaves the block to be transferred “in the background” without CPU involvement
  4. An IRQ typically signals that the operation has completed
    - The data requested is now in the main memory buffer

# Steps to Initiate DMA transfer

- The CPU:
  - Programs the DMA controller for read or write
  - Loads the address register with the start address of the transfer
  - Loads the data count register with the number of bytes to be transferred
- **Note:** At this point the CPU will no longer be directly involved as the DMA controller will communicate directly with memory for the data transfer

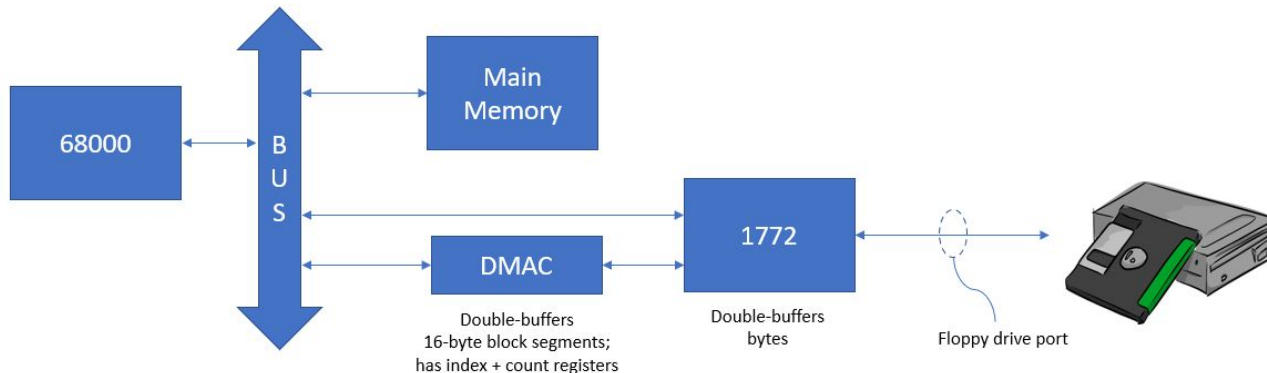

# Steps to Initiate DMA transfer

- The DMA controller periodically requests the bus. When the bus is granted, the DMA controller uses it to transfer data.
  - The controller may use “cycle stealing” to transfer a small number of bytes at a time, or, if it has larger buffers, it may transfer the data in bursts.

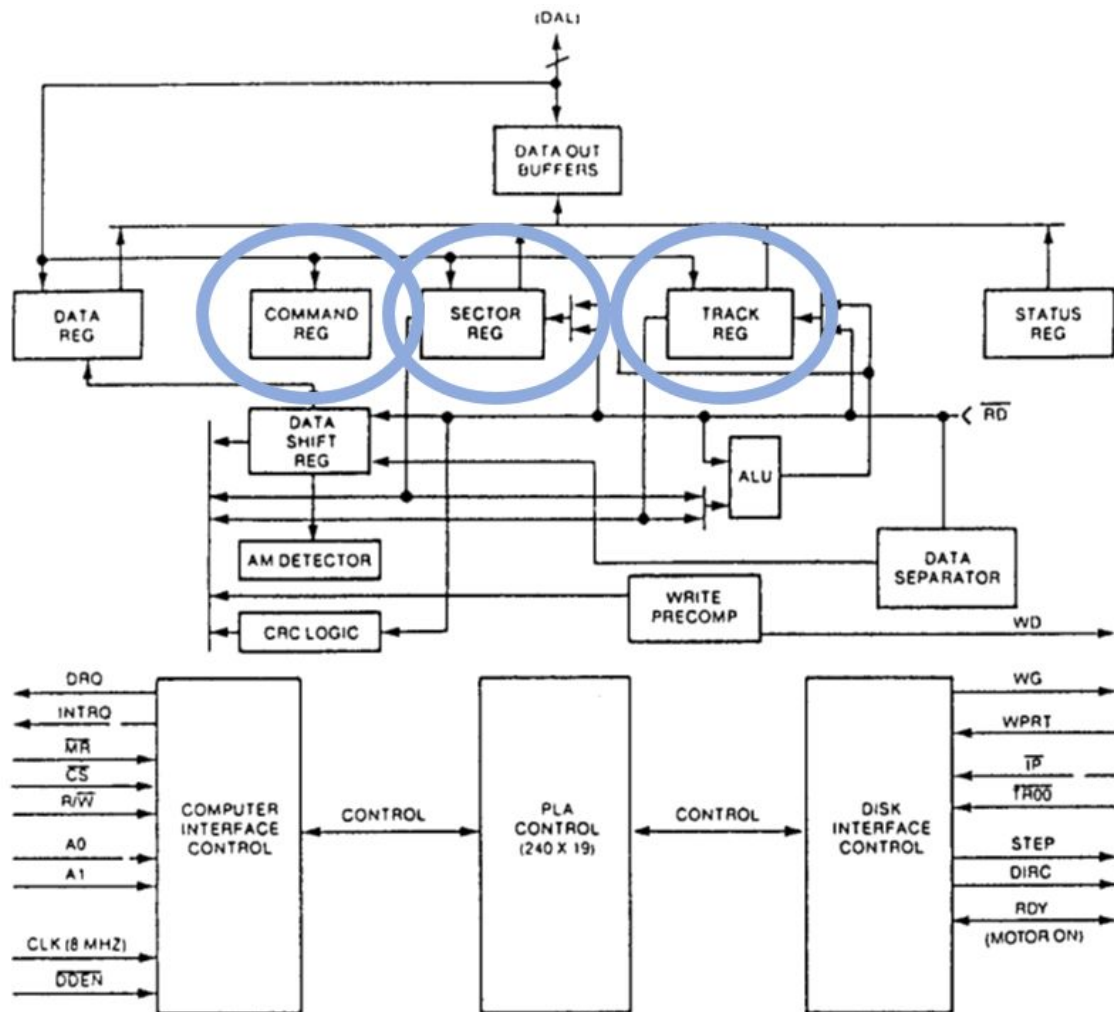


# Simple Example - FDC on Atari ST

- Floppy drive controller:
  - The WD1772
  - Connected to system bus via the DMA controller

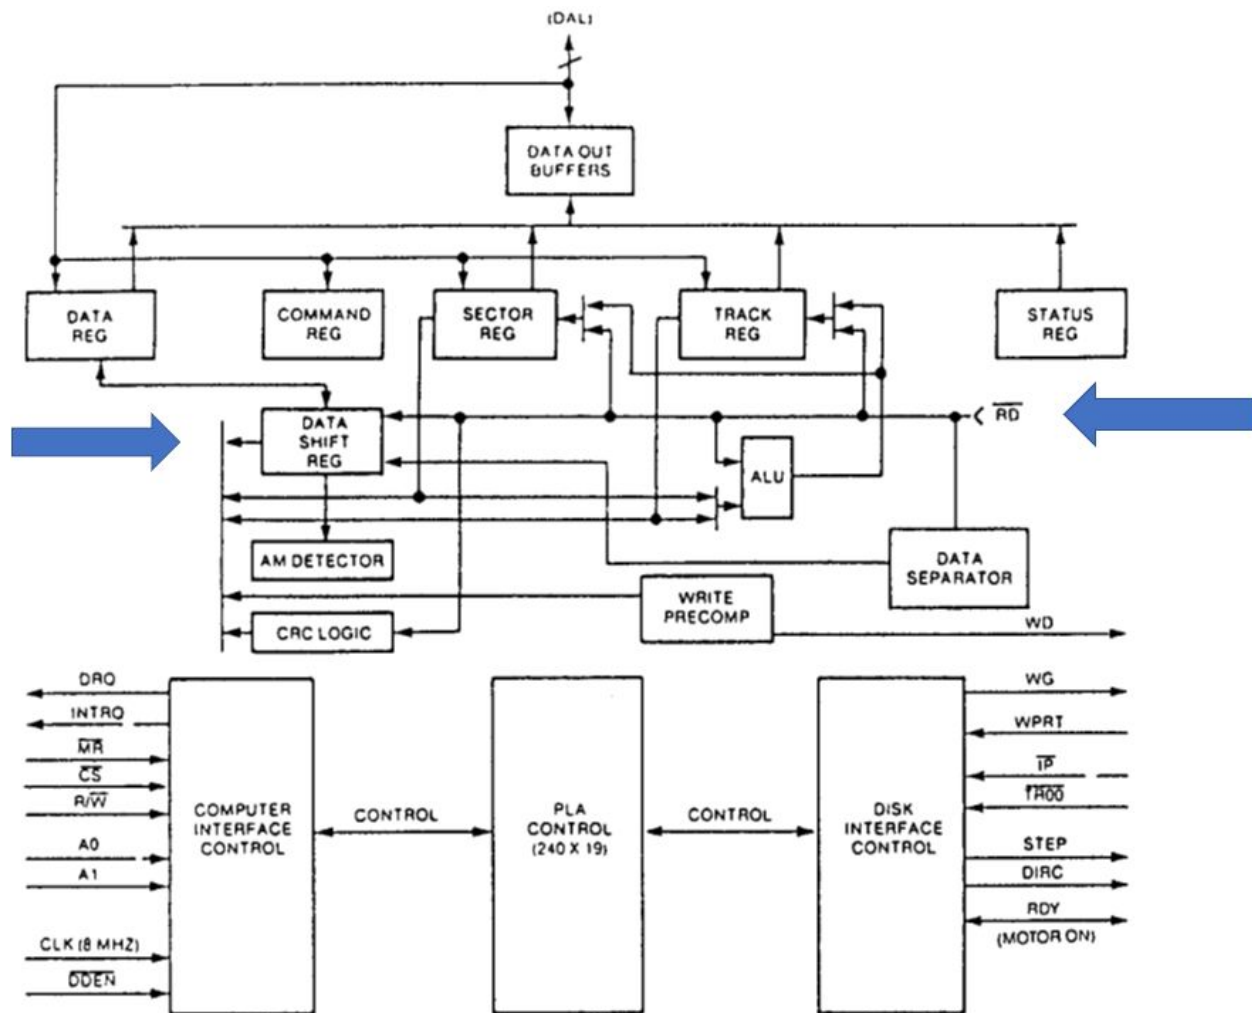


Block  
diagram of  
1772  
internals



Initiate a read  
command via  
write to  
command  
register

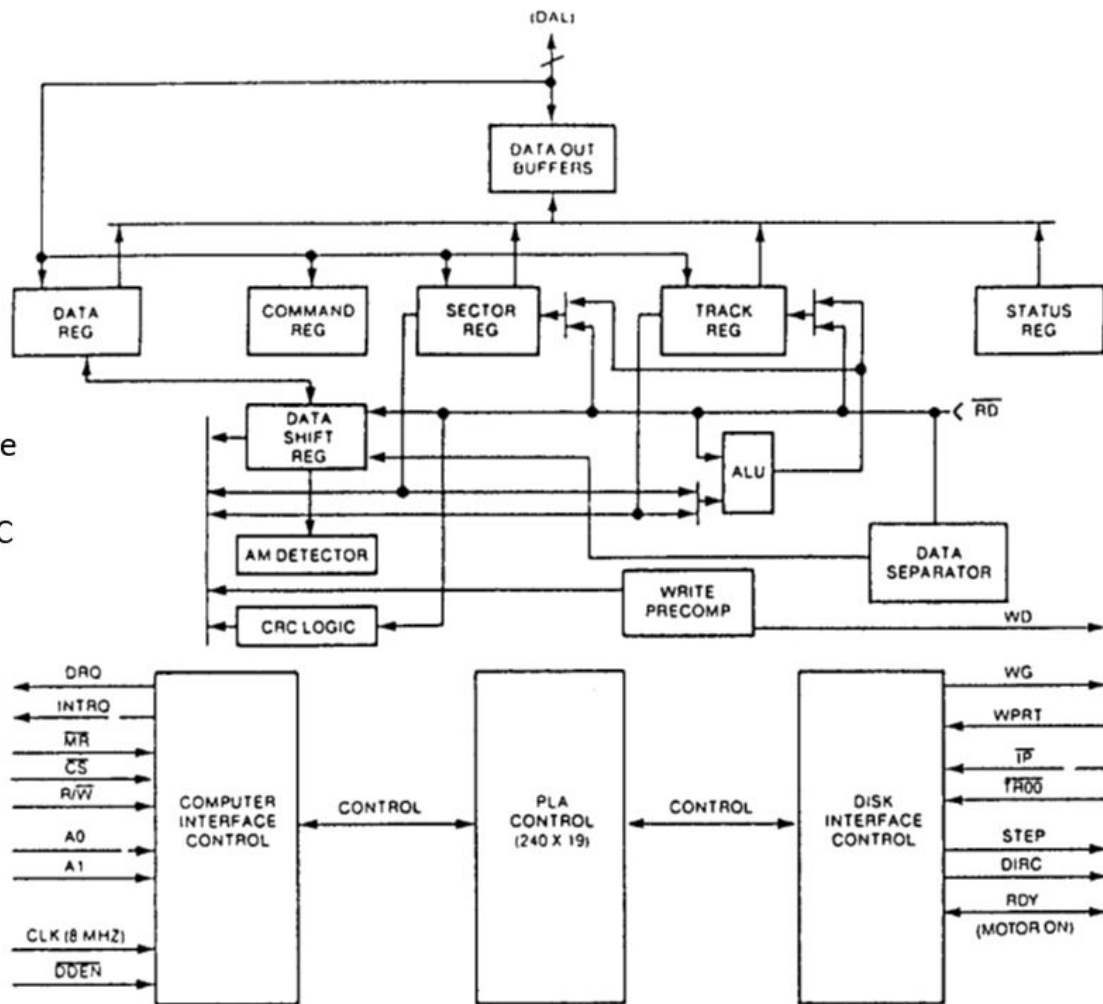
(Block location  
on disk  
specified via  
track/sector  
registers)



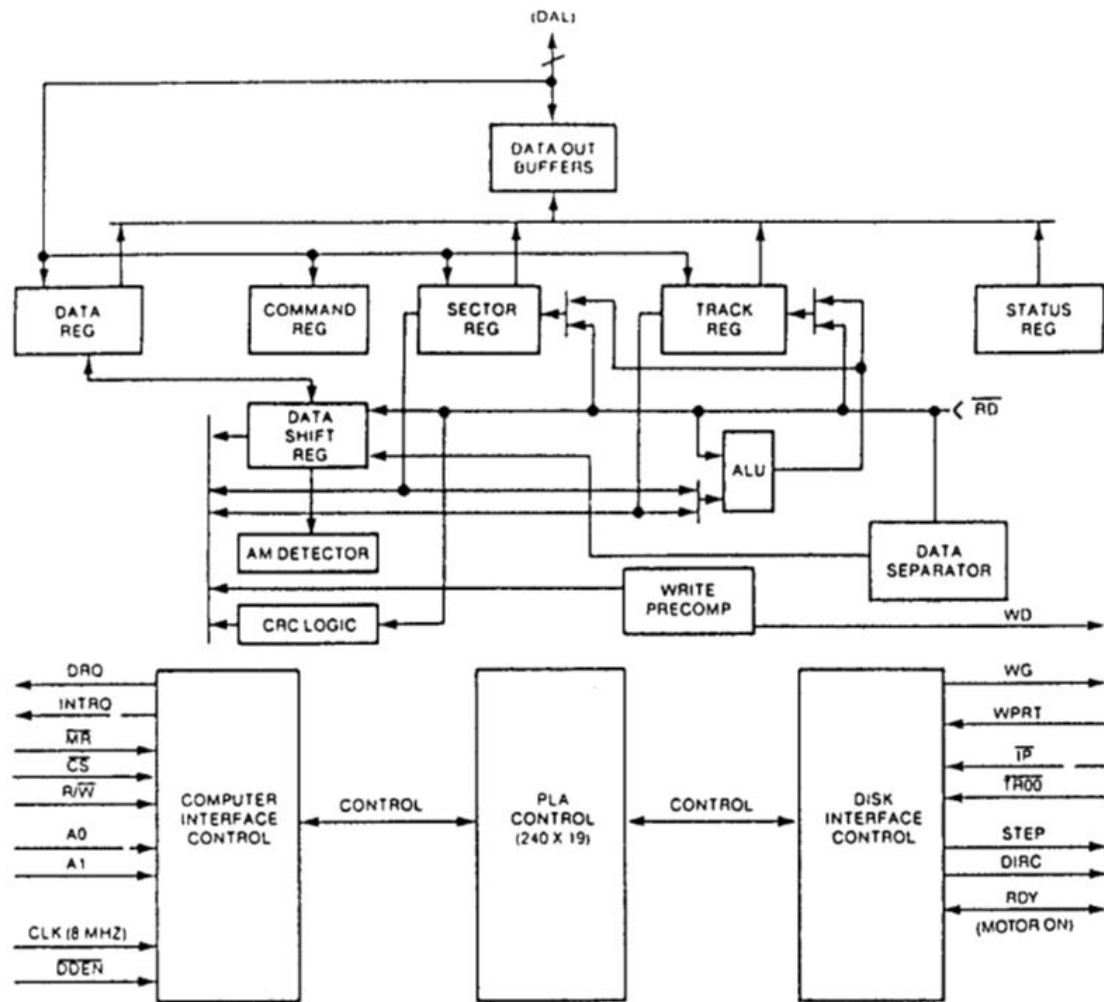
Bytes arrive  
serially over  
floppy port

Fully assembled byte is ready to be taken by DMAC

1772 signals "data request" to DMAC

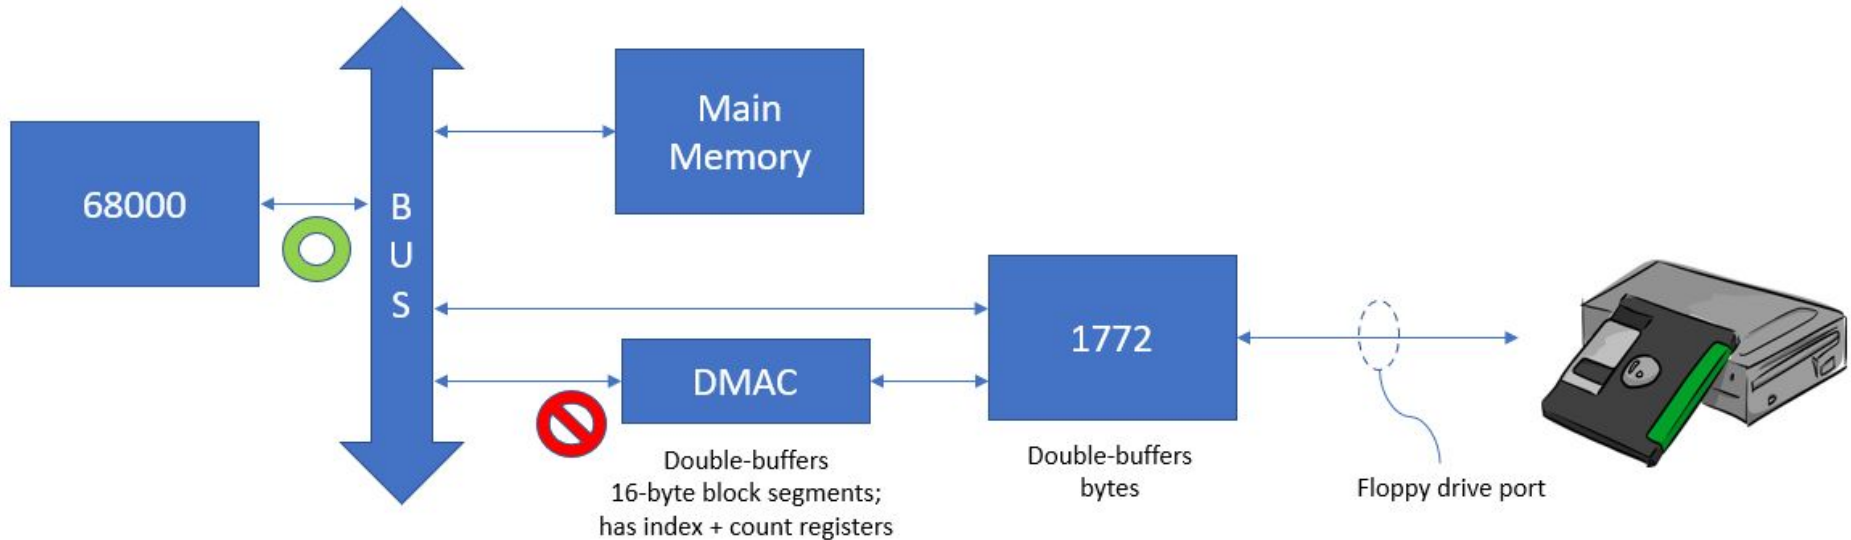


Read by  
DMAC



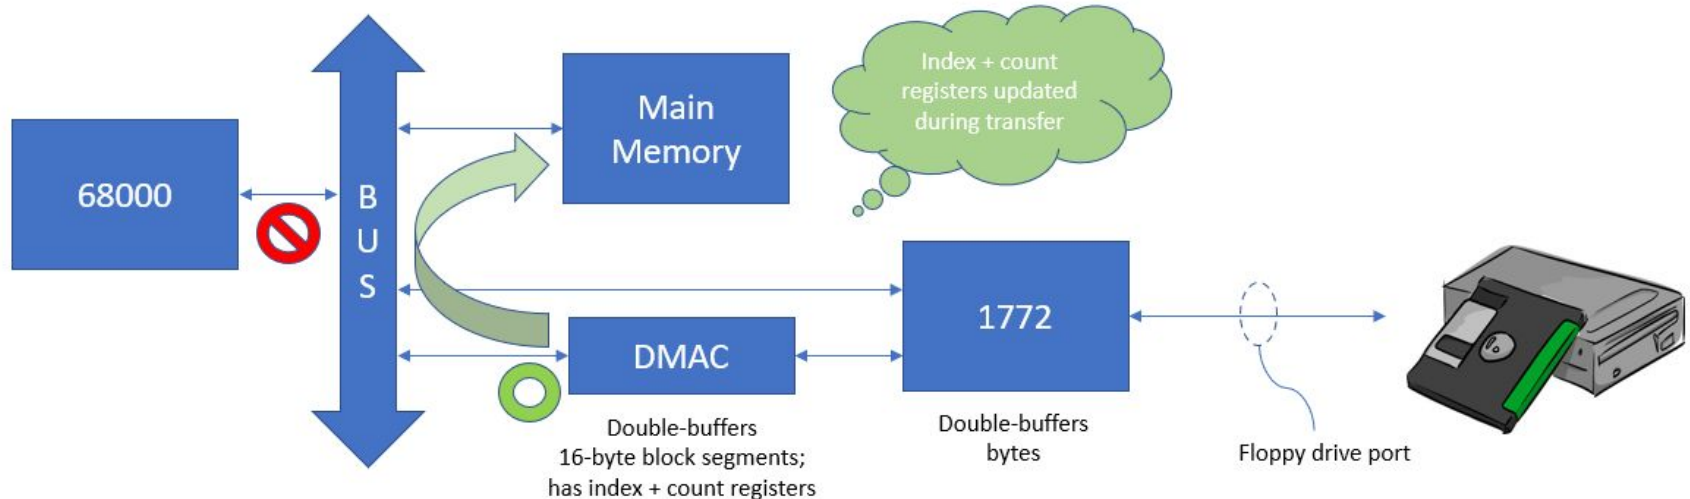
# Simple Example - Atari ST's DMAC

- By default, the CPU is the "bus master"



# Simple Example - Atari ST's DMAC

- DMAC periodically requests to take control of the system bus
- To transfer the next segment of the incoming block directly to memory



# DMA Interfaces

- An I/O interface may have its own built-in DMA controller
- Alternatively, a distinct DMA controller may manage a set of I/O interfaces, which may not themselves be connected directly to the bus
- E.g. the 68000 supports DMA via its BR, BG and BGACK signals
- E.g. the ST has a DMA controller which manages the hard drive and the floppy drive(s)
  - Actually, the memory indexing part of the DMA functionality is performed by the MMU chip.



# Bus Mastering

- The bus master (here, 68000) supports a handshaking protocol with the DMAC:
  - **Bus Request** ← from DMAC to 68000
  - **Bus Grant** ← response from 68000 (when willing)
  - **Bus Grant Acknowledge** ← asserted by DMAC until transfer complete

# Bus Mastering

- DMACs typically transfer data in one of two modes:
  - Burst
    - Upon being granted the bus, all buffered bytes are transferred one after the other
  - Cycle stealing
    - One value is transferred per bus grant
    - Allows interleaved use of bus by CPU and DMAC
    - Cycle-stealing can be used by a DMA-enabled video controller to continually scan a frame buffer as it generates a video signal

# Final Notes

- Also to add:
  - IO interface → IO controller → **IO (co)processor (e.g., GPU)**
    - Start with basic IO interface providing little more than buffering and polling support
    - ... → add interrupt capability (and PIC support)
    - ... → add storage & transfer capability (maybe with separate DMAC support)
    - ... → **add compute capability (+ more onboard memory and/or DMA capability)**