
Serial I/O


— One thing after another —

Overview

- An n -bit value can be transmitted serially, meaning one bit at a time (over n clock cycles).
- The alternative is to transmit in parallel (discussed and compared later).
- Serial communication is accomplished using Transmit (Tx) and Receive (Rx) shift registers.

Overview

- E.g. illustrate transmission of 4-bit value assuming a shared clock signal.



- The rate of serial transmission is measured in bps (bits per second) – and often in kbps and Mbps.

Overview

- Exercise: assuming a serial connection between a keyboard and a computer, what would be an appropriate minimum transfer rate?
- What is the average typing speed?
- Average of 50 – 70 wpm (words per minute):
- With an average of 5 characters per word this is:
- $(70 * 5) / 60 \text{ sec} = 350 / 60 = 6 \text{ characters / sec}$
- Since each character has ~ 8 bits / character
 - this is $6 * 8 = 48 \text{ bits / sec}$
- See http://en.wikipedia.org/wiki/Words_per_minute

Terms

- Baud rate
 - Is a signalling rate (i.e. number of symbols transmitted per second)
 - Not the same as *data rate* (Except if the only signals are 0, and 1)
 - Baud rate is more general than data rate
- Duplex
 - Simultaneous bidirectional communication
 - Requires two data lines

Terms


- Half-Duplex
 - Non-simultaneous bidirectional communication
 - Requires one data line
- Simplex
 - Unidirectional communication
 - Requires one data line

Synchronous Transmission


- Serial communication may be synchronous or asynchronous
- In synchronous, transmission between sender and receiver is synchronized using a common clock
- In synchronous communications, bits are sent continuously
 - If the sender has nothing meaningful to send, bits are still sent
 - There must be a communications protocol so that the receiver can distinguish meaningful data from “pad” data
 - The protocol may include transmission of synchronization or resynchronization data

Synchronous Transmission

- One possible way to synchronize clock over transmission:



- Another possible way to synchronize clock over transmission:



Synchronous Transmission

- How can the clock signal be recovered from a data stream such as:



- Idea: the timing of the transitions in the data stream reveal the sender's clock signal sufficiently well
 - Assuming transitions happen often enough
- A full discussion of synchronous communication protocols is beyond the scope of this course

Asynchronous Communication

- In asynchronous serial communication, there is no common clock to synchronize sender and receiver
- Communication is not continuous, but starts and stops
- The start time of a new “frame” of data is not arranged in advance
- When data is not being transmitted, the line is left in an idle state
 - Either low or high but high is most commonly used
- The frames are normally kept short
 - 5 to 9 data bits – 8 is the most common


Asynchronous Communication

- Each frame has the following format:
 - Start bit ← non-idle level to signal start of frame
 - Data bits ← usually LSB first
 - Optional parity bit
 - 1, 1 ½ or 2 stop bits ← return to idle level (minimum duration)
- Exercise:** assuming high bit for idle, 8 data bits LSB first, no parity bit and 1 stop bit (8N1), what is the value being transmitted below?



Asynchronous Communication

- Labelling the transmission allows for identification of the key elements in the frame:



- First 0, is the start bit (i.e. no longer in idle, high, transmission)
- First bit after the start bit is the LSB
- Next 8 time segments after start bit are the data
 - Data transmitted: 0110 0001 or \$61 or 97_{10} ('a')

Error Detection - Parity


- A full discussion on error detection in transmission is beyond the scope of this course
- Parity uses the number of set bits (i.e. a value of 1) in the original value to allow the receiver to determine if any of changed
 - It is limited to showing if an odd number of bits have changed
 - It is limited so not really used now
 - Does not allow for recovering from a transmission error
 - If parity is being used the parity bit is always inserted

Error Detection - Parity

- “Odd parity” means the sum of the number of set data bits + the parity bit must be odd
 - If the number of 1’s in the data is even the parity bit will be set to 1
 - If the number of 1’s in data is odd the parity bit will be set to 0
- “Even parity” is similar, but the total number of 1’s plus the parity bit is made to be even
 - If the number of 1’s in the data is even the parity bit will be set to 0
 - If the number of 1’s in data is odd the parity bit will be set to 1
- If the receiver detects a parity mismatch, then an error occurred (e.g. due to electromagnetic interference). If not, an error might have occurred.

Error Detection - Parity

- **Exercise:** assuming 7E1, draw a diagram to represent the transmission of ASCII 'T'
 - 'T' == 54 = 0101 0100 (via reference card)
 - Since 7 data bits it becomes 101 0100 as the leading 0 is dropped
 - Using even parity there needs to be an even number of 1's in the data, including the parity bit
 - Since this data has an odd number of data bits the parity bit will be 1



Transmission Details

- Without a common clock, how does the receiver know when one bit ends and the next begins?
- The receiver uses its own high frequency clock to sample the data stream
 - e.g. at $16\times$ the agreed-upon baud rate

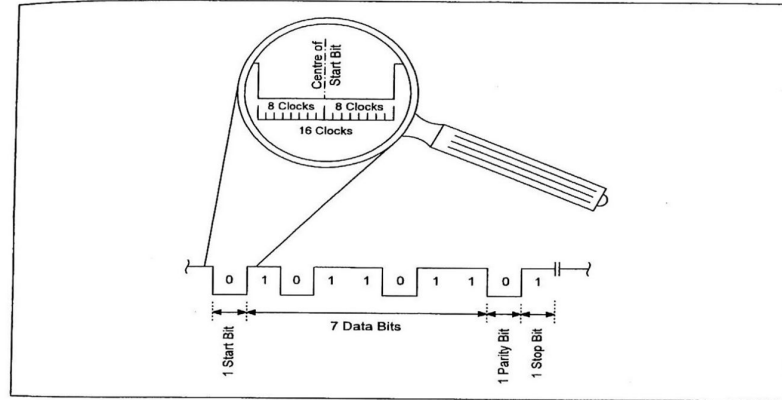


Figure 32.22: Detecting the start bit. The bits are sampled at a frequency 16 times the baud rate to enhance the reliability of data detection.

Transmission Details

- It can therefore detect the start time of a frame (idle to non-idle transition) within a margin of $1/16$ th the transmitter clock period
- It can then attempt to sample the value of the first data bit in the middle of it's transmission period - $16 + 8$ clock ticks later
- Subsequent data bits are sampled similarly - every 16 clock ticks
- The sender and receiver clock are not in sync
 - They will drift Error will accumulate at each sample
 - As long as frames are short, the error will not accumulate to the point where bits are being sampled outside their beginning and end times

Transmission Details

- The receiver will re-sync its clock on the start bit




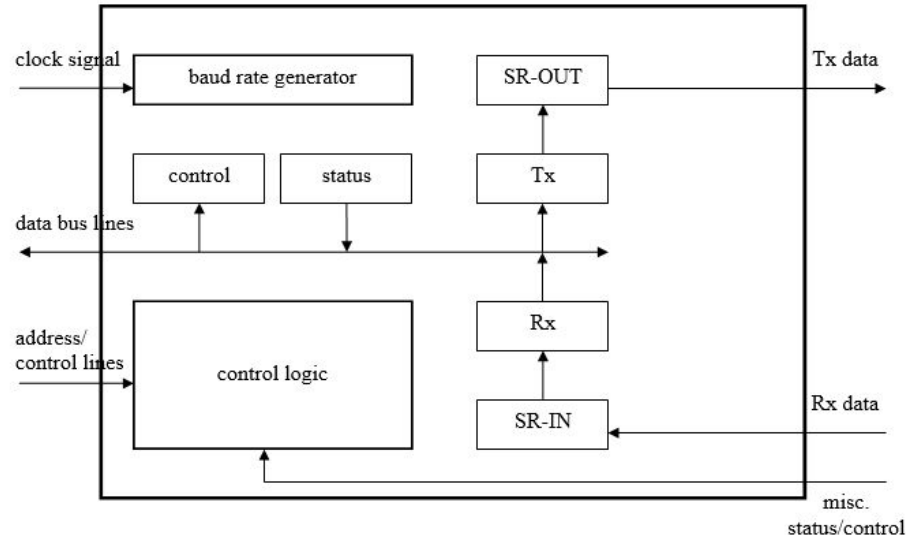
Figure 6.42 Relative timing of data and receiver clock in an asynchronous system.

Simple UART

- A UART (Universal Asynchronous Receiver Transmitter) is an I/O interface for asynchronous serial communications
- “Universal” means that the baud rate is programmable, not fixed
 - There may be a command register whose contents specify by how much to frequency-divide a master clock.
- Typical CPU-accessible registers:
 - Command (control)
 - Status
 - Tx
 - Rx

Simple UART

- Generalized organization and external signals:
 - Note the double buffering (e.g. for transmission, the Tx and SR-OUT registers)



Simple UART - Control Register Bits

- Master clock divide-down for baud rate generation
- Frame format settings (e.g. 8N1 vs. 7E2)
- Rx interrupt enable
- Tx interrupt enable

Simple UART - Status Register Bits

- Tx empty
- Rx full
- Rx overrun (error)
- Parity error
- Framing error ← received stop bit not valid
- Interrupt requested

A Simple UART

- A USART is a similar type of interface, capable of both synchronous and asynchronous operation
- ACIA (Asynchronous Communications Interface Adapter) is another name for a UART
 - The Atari ST has two MC6850 ACIAs: one for managing the serial connection to the keyboard, the other for managing the MIDI port
 - It also has a USART onboard its 68901 chip (discussed later).
- The next tutorial completes the discussion of the MC6850

Serial Line Codes

- Above, it was assumed that:
 - To transmit a 0, a low level is placed on the line
 - To transmit a 1, a high level is placed on the line
- This is the most obvious serial line code, called NRZ (Non Return to Zero)
- There are others, which are sometimes preferable – esp. for synchronous communication
- Why?

Serial Line Codes


- A digital phase-locked loop (DPLL) is an analog/digital circuit that can be used to recover a clock signal from a serial data stream
 - It only works if there are enough 0-1 or 1-0 transitions in the data to reveal the clock signal
- Using NRZ, a DPLL will fail if too long a stream of 0s or 1s is transmitted
- Why would such a stream (one with a long series of one data value) occur?
 - An image file

Serial Line Codes

- Another reason: some storage media (e.g. magnetic disks) don't store a bits as absolute 0s and 1s
 - They store the transitions between bit levels. So, NRZ is unsuitable.
- There are a few other reasons why alternatives to NRZ are sometimes used, but they are beyond the scope of this course
 - Example: media characteristics, electrical reasons – distance, repeating
- A full discussion of serial line codes is beyond our scope

Serial Line Codes

- The following chart illustrates some common ones:
 - The point here is to understand that there are many ways to transmit binary values



Invert transmission on a 1
This is used by USB

Return to zero

Zero is 0 signal
1 alternates between ve
and -ve

1 is 01 transition
0 is 10 transition

NRZI = NRZ invert on 1s
BPRZ = bipolar RZ

Serial Line Codes - Tradeoffs

- Some are better for clock recovery than others
 - More transitions
- Some are suitable for transition-sensitive media storage
- Some are “balanced”
 - There is no DC component to bias a receiving circuit
- Some require higher bandwidth due to high frequencies
 - Lots of fast transitions
- See: <http://www.interfacebus.com/Definitions.html>

RS-232

- RS-232 is a standard for the serial interconnection of two components:
 - a DTE (Data Terminal Equipment) e.g. a computer
 - a DCE (Data Communications Equipment) e.g. a modem
- RS-232 defines a physical interface and a set of signals.



RS-232

- Nine of the signals are commonly used:
- G common Ground
- TD Transmit Data generated by DTE
- RD Receive Data
- DTR Data Terminal Ready generated by DTE
- DSR Data Set Ready
- RTS Request To Send generated by DTE
- CTS Clear To Send
- DCD Data Carrier Detect
- RI Ring Indicator

RS-232

- A DE-9 port is common, although other variants (including a DB-25 connector) are also allowed.
- D is the shape of the metal shield and B-E are the size
- Although there are several serial communications standards, the term “serial port” usually refers to a 9- or 25-pin RS-232 port
- See <http://en.wikipedia.org/wiki/RS-232>
- See <https://en.wikipedia.org/wiki/D-subminiature>

DSubminiature Plug Ends



RS-232

- “Null modem” cables can be used to directly connect two computers via their serial ports
- Once very common, RS-232 ports have almost disappeared from personal computers over the past few years
 - In favour of USB, etc
- RS-232 is still common in several settings
 - Communicating with microcontrollers, etc
- UART chips are often designed to be RS-232 compatible
 - i.e. with RS-232 specific pins and functionality

RS-232

- This was “recently” found at store in Calgary when the store opened
- This was printed out as a result of turning on the checkout.
- Cash registers use an RS-232 connection

DresserWayne


DW-12 Ver.CU31-0200

```
Serial Interface
Baud rate      : 115200bps
Data bit       : 8 bits
Parity         : None
Flow Control   : XON/OFF
I/F Level      : RS-232C
```

```

Dip switches
      DS1
    1 2 3 4 5 6 7 8
ON   0 0
OFF  0 0 0 0

```

!"#\$%&'()*+,-./01234567
89:;<=>?@ABCDEF GHIJKLMNO
PQRSTUVWXYZ[\]^_`abcdefg
hijklmnopqrstuvwxyz{|}~
CUEAAAAACBEEIILIAAEZAZOOUU
YUUCEFYHfaiouüñnao{~}!«»

א ב ג ד ה ו ז ח ט י כ ל מ נ ס ע פ צ ק ר ש ת

USB

- USB is a newer standard for connecting peripherals
 - The latest version is USB 3.1
- On PCs, USB is now used for almost everything except internal drives, graphics and specialty PCI expansion cards
- Compared to RS-232, USB is a very sophisticated standard. A complete study is beyond the scope of this course

USB - Goals of the Specification

- Flexibility and “ease of use”
 - Including multi-peripheral connections, use for user not programmer
- Low cost
- Small, robust connectors
- High speed
 - Fast enough for external storage devices and real-time audio, video and voice
- Dynamic attach / detach
- Bandwidth and latency guarantees

USB


- A USB system is comprised of:
 - A USB host
 - USB devices
 - The USB interconnect
- The USB host is the host computer. The computer has a USB controller – this is the I/O interface to the USB system (the “USB bus”)
- Actually, even though several USB devices can be serviced through one USB controller, PCs often ship with 4-5 controllers

USB

- A USB device is one of:
 - A USB hub, which provides USB device connection points
 - A USB function (e.g. a keyboard, joystick, memory stick, digital camera, ...)
- The USB controller has an integrated “root hub”. Therefore, a USB system is arranged as a tree (a “tiered star bus topology”)
- The USB interconnect is the manner in which USB devices are connected to and communicate with the host

USB Connectors

- USB uses 4-pin connectors
 - type A and B shown below
 - there are also miniature connectors
- Two of the pins carry power
- The other two carry serial data using half-duplex differential signaling
- The cable is shielded, and the data lines are twisted pair



USB

- Data is transmitted serially using NRZI (invert on 0s, with bit-stuffing after five consecutive 1s to ensure sufficient transitions).
 - i.e. communication is synchronous, with clock recovery from the data stream
- The USB bus protocol is sophisticated. A few highlights:
 - Detection of device attach and detach
 - Negotiation for power requirements
 - Negotiation for bandwidth and latency (all communication is scheduled by the host controller)
 - Error detection and self-recovery for lost packets

USB

- The bus is a “polled” bus – the USB controller schedules and initiates all data transfers (again, half-duplex).
- The USB controller repeatedly enumerates the connected devices. Devices are required to be able to identify themselves when attached
- Each device has a “device class”. The set of classes is standard. An O/S is required to provide default drivers for all classes, so that the appropriate driver can be loaded when a device is attached. There is a reserved device class for unique devices which require custom drivers

USB

- An O/S provides a USB controller driver, plus drivers for each device class
 - For ease of use, USB libraries are typically used to provide an additional layer of abstraction.
- USB 2.0 supports a data rate of 480 Mbps!
 - Note: This does not mean that the throughput of an individual device is 480 Mbps
 - The bus is shared with other devices and there is protocol overhead.
 - USB 3.0 supports 5.0Gbps, 3.1 10 Gbps, 3.2 20 Gbps USB4 40 Gbps and v2.0 80Gbps!

USB

- Specs: <https://en.wikipedia.org/wiki/USB>
- How USB works: <https://youtu.be/wdgULBpRoXk>